

Exploratory Data Analysis for Smart Factory Predictive Maintenance Using NASA Turbofan Engine Sensor Data

Module: IOT106TC Big Data Analytics

Assignment: Assignment 1 — Exploratory Data Analysis Report

Scenario: Scenario 1 — Smart Factory Predictive Maintenance

Dataset: NASA C-MAPSS Turbofan Engine Degradation Simulation — FDOO1 Subset

Companion notebook: notebooks/Assignment1_SmartFactory_EDA.ipynb

Executive Summary

This report presents the Exploratory Data Analysis (EDA) of the NASA C-MAPSS Turbofan Engine Degradation Simulation dataset (FDOO1 training subset), in a Smart Factory predictive maintenance framing. The file `train_FDOO1.txt` comprises **100** engines $\times$ **20,631** cycle-level rows and **26** raw sensor/setting columns—21 anonymised IoT sensor channels plus three operational settings. The notebook calculates Remaining Useful Life (RUL) from run-to-failure labels, derives lifecycle features (`cycle_ratio`, `life_stage`, `near_failure`, rolling means), and screens all sensors. Approximately **seven** sensor channels exhibit near-zero variance (threshold: standard deviation < 0.01) and contribute little diagnostic information. **Fourteen** non-constant channels show $|\text{Pearson } r| > 0.3$ with RUL; the strongest single relationship is **sensor_11** ($r = -0.696$). **15.0%** of rows fall in the exploratory near-failure band ($\text{RUL} \leq 30$ cycles). The cleaned dataset `data/processed/SmartFactory_cleaned.csv` contains **42** columns and is ready for Assignment 2 RUL regression and near-failure classification. Figures 1–10 below are produced by the notebook and saved under `figures/`.

1. Introduction and Business Understanding

1.1 Business Context

Modern industrial environments increasingly rely on IoT sensor networks to monitor the health of critical machinery in real time. In aviation and manufacturing, turbine engines and turbofan units represent high-value, safety-critical assets where unplanned failure carries severe operational, financial, and safety consequences. The traditional maintenance approach — fixed-interval scheduled maintenance — results in either premature replacement of components that still have useful life remaining, or reactive maintenance following unexpected failures. Both outcomes are economically inefficient and, in the latter case, potentially hazardous.

Predictive maintenance addresses this gap by analysing continuous sensor data to estimate how much useful life an asset has remaining, enabling maintenance to be scheduled at the optimal point: early enough to prevent failure, late enough to avoid unnecessary servicing. This Assignment uses the NASA C-MAPSS FDOO1 dataset as a proxy for an industrial IoT scenario in which sensor-equipped engines are monitored from installation to end-of-life.

1.2 Problem Statement

The central problem for Assignment 1 is not yet predictive modelling, but **data understanding and preparation**. Before a reliable RUL prediction system can be built, three questions must be answered:

1. **What does the data look like?** Structure, quality, completeness, and scale.
2. **What patterns are present?** How do sensors behave over an engine's lifecycle? Which channels carry useful degradation signals?
3. **What features should be engineered?** Which representations will support Assignment 2 modelling?

The goal is to produce a clean, well-understood dataset and a grounded set of findings that inform both maintenance operations and future modelling strategy.

1.3 Stakeholder Analysis

Stakeholder	Primary Concern	Relevance to This Analysis
Plant / Fleet Manager	Minimise unplanned downtime	Understanding which engines are near failure and why
Maintenance Director	Schedule preventive maintenance efficiently	Identifying degradation onset timing across the fleet
Finance Director	Reduce cost of unnecessary maintenance and downtime losses	Quantifying sensor signal quality to prioritise investment in data-driven maintenance
Safety / Reliability Team	Prevent catastrophic failure	Identifying near-failure engine states and most informative warning signals
Data / IoT Engineering Team	Maintain sensor pipeline; build predictive models	Understanding which sensor channels to prioritise; selecting features for Assignment 2

1.4 Success Criteria

The analysis is considered successful if:

- RUL is correctly calculated and validated for all engines.
- Data quality is fully assessed, documented, and addressed.
- At least 8 engines are analysed and compared.
- At least 10 sensor channels are explored in depth.
- Informative sensor channels are identified and ranked.
- Engine degradation patterns are characterised and interpreted.
- A cleaned, feature-enriched dataset is produced and saved.
- Business-relevant recommendations are supported by evidence.
- Responsible AI use is documented in the AI Collaboration Report.

2. Data Understanding

2.1 Data Source and Structure

The dataset used in this analysis is the **NASA C-MAPSS (Commercial Modular Aero-Propulsion System Simulation) Turbofan Engine Degradation Simulation Data**, specifically the **FD001 training subset**. The data was generated by NASA's Prognostics Center of Excellence (PCoE) and is widely used in the predictive maintenance research community as a benchmark for RUL prediction algorithms.

Dataset source: NASA PCoE Data Repository

File used: train_FD001.txt

Operating condition: Single flight condition (FD001 operates under one fixed operational regime, making normalisation straightforward)

Failure mode: High Pressure Compressor (HPC) degradation (single fault mode)

The dataset is a **run-to-failure** dataset: each engine is recorded from its first operational cycle until it fails. The last recorded cycle for each engine is its failure cycle. This design allows Remaining Useful Life to be unambiguously calculated in retrospect.

2.2 Dataset Schema

The file contains no header. Data is whitespace-separated with 26 columns, assigned as follows:

Column	Name	Description
1	unit_number	Engine identifier
2	time_in_cycles	Operational cycle index (1 = first cycle)
3–5	operational_setting_1/2/3	Flight condition parameters
6–26	sensor_1 to sensor_21	IoT sensor channel readings

Sensor physical identities are not explicitly documented in the dataset file. To avoid overclaiming physical meaning, this analysis refers to all channels as anonymised sensor labels (sensor_1 through sensor_21).

2.3 Data Quality Assessment

Quality Metric	Value
Total rows (raw)	20,631
Total rows (after deduplication)	20,631
Columns (raw schema)	26
Missing values (total)	0
Duplicate rows removed	0
Near-constant sensor channels (std < 0.01)	7
Rows with ≥ 1 sensor value outside the $3 \times \text{IQR}$ fence (<i>flag only; no row drops</i>)	1,911
Engineered features added (see notebook)	16

Missing values: None detected in train_FD001.txt across all loaded columns (outputs/data_quality_summary.csv).

Duplicate rows: Zero exact duplicates; row count unchanged.

Near-constant sensors: Channels **sensor_1**, **sensor_5**, **sensor_6**, **sensor_10**, **sensor_16**, **sensor_18**, and **sensor_19** all have pooled standard deviation **below 0.01** under the notebook rule. Correlation with RUL is undefined for strictly constant series and negligible for the marginal-variance cases. These channels are **retained** in the exported CSV but **marked as low-information** for monitoring prioritisation (outputs/sensor_rul_correlation.csv).

Outlier handling policy: Row-level IQR fencing (multiplier $3 \times$, **any sensor**) **flags 1,911 rows** once per row if any channel breaches its fence—**without row deletion**. Late-life extremes are treated as plausible degradation artefacts.

Figure 1 shows the pooled standard deviation ladder (with the 0.01 rule-of-thumb cutoff) alongside the absence of missing readings.

Figure 1: Data Quality Summary




Figure 1. Data-quality summary plots generated in the notebook.

2.4 Initial Statistical Summary

Engine-level lifetime `max_cycle` per `unit_number` ranges from **128** to **362** cycles (**mean ~206, median 199, standard deviation ~46**). The positively skewed distribution (**skew ≈ 1.02**) indicates a long tail of **long-lived** engines rather than symmetry around the mean—a pattern consistent with heterogeneous wear accrual in a stochastic degradation process.

Across all rows, `time_in_cycles` spans **1–362** (**mean ~108.8, std ~68.9**). Uncapped `RUL` mirrors the same dispersion (**max 361**) because `RUL = max_cycle - time_in_cycles`.

Illustrative univariate dispersion for `sensor_2` (among the trending channels used in lifecycle illustrations): approximately **641.21–644.53** (**mean ~642.68, std ~0.50**).

Operational settings `operational_setting_1` and `operational_setting_2` show **minimal** dispersion (standard deviations ~0.0022 and ~0.00029; only a handful of unique values). `operational_setting_3` is **constant at 100.0** for every observation—consistent with FDO01's single-flight-regime textbook description.

3. Data Cleaning and Preparation

3.1 Cleaning Strategy and Rationale

The cleaning strategy follows a conservative, domain-aware approach that prioritises **preserving signal over removing noise**:

1. **Missing values:** Column-wise scan—none found on FDO01 train.
2. **Duplicate rows:** No exact duplicates removed—count remained zero throughout.
3. **Constant features:** Flag channels with pooled std < 0.01; **do not drop** preemptively ahead of Assignment 2 feature selection.
4. **Outlier screening:** Flag using a $3\times\text{IQR}$ fence; **do not auto-delete**.
5. **Type validation:** All 26 operational/sensor inputs are numeric; no coercion failures.

3.2 RUL Calculation Methodology

RUL at cycle t was calculated as:

$$\text{RUL}_t = \text{max_cycle_for_that_engine} - \text{time_in_cycles}_t$$

Here `time_in_cyclest` is the operational cycle index for that row, and `max_cycle_for_that_engine` is that engine's final recorded cycle. For each engine, the notebook verified that the final recorded row had `RUL = 0`. All **100** engines passed this check.

3.3 Feature Engineering

Beyond the raw columns, the engineered fields include `max_cycle`, `RUL`, `cycle_ratio`, ordinal `life_stage` (early: `cycle_ratio` ≤ 0.33 ; mid: ≤ 0.67 ; late: else), exploratory `near_failure` (indicator for `RUL` ≤ 30), `capped_RUL`, and `sensor*_ro1120` rolling means for the **top 10 absolute RUL-correlated** active sensors (**window = 20** cycles within each engine timeline). See `outputs/feature_summary.csv` for the definitive column glossary.

3.4 Before-and-After Data Quality Report

Metric	Before	After (cleaned CSV)
Total rows	20,631	20,631
Columns	26	42
Missing values	0	0
Duplicate rows	0	0
Near-constant sensors flagged	—	7
Rows flagged by 3×IQR rule (≥ 1 sensor)	—	1,911
Engineered numeric/categorical derivatives	0	16

4. Exploratory Data Analysis

4.1 Engine Lifecycle and RUL Distribution

The training split tracks **100** independent engines (**20,631** pooled observations).

Lifetime distribution: `max_cycle` min **128**, max **362**, mean $\approx$ **206.31**, median **199**, $\sigma \approx$ **46.34** (Figure 2). **Tertile cut-points** on `max_cycle` separate **short** ($\leq$ **188** cycles), **medium** (**188–213** cycles), and **long** ($>$ **213** cycles) lifetime cohorts used for comparative plots.

RUL distribution: Per-row `RUL` $\in$ [**0**, **361**]; the histogram is **right-skewed** (Figure 3) because early-life rows contribute many high-RUL samples. **15.0%** of rows satisfy **RUL** $\leq$ **30** (`near_failure` = 1). Outside that band, mean `RUL` is $\approx$ **124** cycles (median **118**), illustrating a long **pre-critical** monitoring window.

Life-stage counts (row-level tertiles on `cycle_ratio`): early **6,767**, mid **6,999**, late **6,865**—balanced by construction.

Figure 4 traces `sensor_2` (strong $|r|$ with `RUL`) for **eight** stratified engines spanning the lifetime spectrum.




Figure 2. Engine lifetime distribution.




Figure 3. RUL and lifecycle-stage composition.




Figure 4. Multi-engine lifecycle examples (≥ 8 engines required).

4.2 Univariate Sensor Analysis

All **21** sensors were screened for variance and correlation with RUL (outputs/sensor_rul_correlation.csv).

Near-constant (7): **sensor_1, sensor_5, sensor_6, sensor_10, sensor_16, sensor_18, sensor_19.**

Active channels: 14 sensors remain after excluding the seven low-variance channels.

Channels with $|r(\text{sensor}, \text{RUL})| > 0.3$: **14** sensors — **sensor_11, sensor_4, sensor_12, sensor_7, sensor_15, sensor_21, sensor_20, sensor_2, sensor_17, sensor_3, sensor_8, sensor_13, sensor_9, sensor_14** (descending by $|r|$).

Top five by |RUL correlation| — empirical ranges on raw train rows:

- **sensor_11:** 46.85–48.53 (mean **47.54**, σ **0.267**).
- **sensor_4:** 1382.25–1441.49 (mean **1408.93**, σ **9.00**).
- **sensor_12:** 518.69–523.38 (mean **521.41**, σ **0.738**).
- **sensor_7:** 549.85–556.06 (mean **553.37**, σ **0.885**).
- **sensor_15:** 8.32–8.58 (mean **8.44**, σ **0.038**).

These statistics support monotonic **trendability** rather than discrete jump behaviour.

4.3 Sensor–RUL Relationships

Figure 5 ranks **Pearson $r(\text{sensor}, \text{RUL})$** for all 21 channels (grey bars mark the seven low-variance sensors). The top five signed correlations are **sensor_11 (–0.696)**, **sensor_4 (–0.679)**, **sensor_12 (+0.672)**, **sensor_7 (+0.657)**, **sensor_15 (–0.643)**.

Figure 6 bins **cycle_ratio** (20 equal-width bins) and plots population means $\pm 95\%$ CI for the **six** highest- $|r|$ **non-constant** sensors used in the notebook (**sensor_11, sensor_4, sensor_12, sensor_7, sensor_15, sensor_21**).

Figure 9 heatmaps pairwise correlations among those six sensors plus **time_in_cycles**, **cycle_ratio**, and **RUL**, highlighting multicollinearity that Assignment 2 should regularise or compress.

Directionality note: Negative r with RUL implies the channel **rises** as failure approaches; positive r implies it **falls**—both are degradation-informative.

Figure 5: Sensor-RUL Correlation Ranking

All 21 Sensor Channels — Correlation with Remaining Useful Life




Figure 5. Sensor–RUL correlation ranking.

Figure 6: Key Sensor Trends Over Engine Lifecycle
(Population Mean $\pm$ 95% CI across all engines)




Figure 6. Key sensor trends over normalised life.

Figure 9: Correlation Heatmap — Key Sensors and Engineered Features




Figure 9. Correlation structure for modelling-relevant variables.

4.4 Comparative Engine Degradation Patterns

Figure 7 compares **early** / **mid** / **late** distributions for the same six sensors—**interquartile expansion** in **late** life is visible for several channels, supporting heterogeneous end-stage wear.

Figure 8 contrasts **short-life** (bottom quartile of `max_cycle`) versus **long-life** (top quartile) engines: **trajectory shapes align qualitatively**, but **levels and slopes diverge**, so **lifetime cohort** carries signal beyond a simple time rescaling.

**Figure 7: Sensor Distribution Across Life Stages
(Early / Mid / Late)**




Figure 7. Life-stage sensor comparison.

**Figure 8: Short-Life vs. Long-Life Engine Sensor Comparison
(Mean $\pm$ 95% CI over cycle_ratio)**




Figure 8. Short-life versus long-life engine cohorts.

4.5 Pattern and Trend Discovery

1. **Divergence onset varies: sensor_4** and **sensor_11** accumulate visible mean drift from **cycle_ratio ≈ 0.2 onward** (Figure 6), whereas **sensor_12** and **sensor_7** stay comparatively flat until the **upper third** of normalised life, then steepen—consistent with staged degradation behaviour in this simulated run-to-failure dataset.
2. **Late-stage variance inflation:** Elevated spreads in **late** boxes (Figure 7) motivate **ensemble** dashboards rather than single-threshold alarms.
3. **Operational settings constrained:** **operational_setting_3** is constant; **_1 / _2** are near piecewise-constant—lifetime spread (**234** cycle swing) therefore **cannot** be explained by observable flight-condition knobs in FD001 (Figure 2 vs. Section 2.4).
4. **Structured low-information sensors:** Near-constant series are **globally flat**, suggesting low diagnostic variation in these anonymised channels rather than an obvious missing-data or sensor-dropout issue.

4.6 Advanced Analysis: PCA Life-Stage Visualisation

Principal components analysis (standardised) used the ten active sensors ranked highest by absolute correlation with RUL in training (sensor_11, sensor_4, sensor_12, sensor_7, sensor_15, sensor_21, sensor_20, sensor_2, sensor_17, sensor_3).

Principal components 1 and 2 jointly explain approximately 78.8% of feature variance, with PC1 explaining about 74.7% and PC2 about 4.2%. Figure 10 is coloured by life stage (early, mid, late) and by RUL, and shows monotonic stratification along PC1, supporting the view that the selected sensor channels form a learnable degradation-related representation for Assignment 2.




Figure 10. PCA of the top-ten correlated sensor manifold.

5. Business Insights and Recommendations

5.1 Key Findings Summary

Finding 1 — Significant engine lifetime variability

Observation: Lifetimes span **128–362** cycles ($\sigma \approx 46$).

Evidence: Figure 2; `outputs/engine_level_summary.csv`

Business implication: Fixed-interval policies mis-allocate overhaul for early failures and defer attention on unusually durable assets.

Finding 2 — Near-failure spans a minority slice of telemetry

Observation: **15.0%** of readings sit in **RUL ≤ 30** ; remainder averages ≈ 124 cycles of headroom.

Evidence: Figure 3; engineered `near_failure` label

Business implication: Condition-based rules can flag wear **before** the critical band if rolling statistics are tracked.

Finding 3 — Seven channels are effectively non-informative

Observation: **sensor_1, sensor_5, sensor_6, sensor_10, sensor_16, sensor_18, sensor_19** fall below the **0.01** std rule.

Evidence: Figure 1; `outputs/data_quality_summary.csv`

Business implication: Edge filtering or slower sampling on these tags may reduce bandwidth, but this should be validated before deployment beyond the FD001 subset.

Finding 4 — Concentrated RUL correlation in a sensor subset

Observation: **sensor_11** leads at $|r| = 0.696$; follow-on leaders include **sensor_4, sensor_12, sensor_7, sensor_15, sensor_21** (all $|r| > 0.63$).

Evidence: Figure 5; `outputs/sensor_rul_correlation.csv`

Business implication: Dashboards and models should centre on this compact set before expanding breadth.

Finding 5 — Monotonic lifecycle separation

Observation: Binned means and **early/mid/late** boxplots shift coherently for the six primary sensors (Figures 6–7; `outputs/sensor_life_stage_shift.csv`).

Business implication: Trend slopes (optionally on rolling means) are credible early-warning features.

Finding 6 — Lifetime quartiles separate average trajectories

Observation: Short- versus long-life cohort means diverge in level/slope while preserving directional similarity (Figure 8).

Evidence: Figure 8; `outputs/engine_level_summary.csv`

Business implication: Fleet segmentation by historical `max_cycle` may aid risk triage when combined with live sensors.

Finding 7 — PCA compresses variance with interpretable PC1

Observation: **PC1+PC2 $\approx 78.8\%$** variance; **life_stage** and **RUL** colour gradients align with **PC1** (Figure 10).

Evidence: Figure 10; notebook PCA cell

Business implication: Latent representations or regularised regression on PC scores are defensible baselines for Assignment 2.

5.2 Data-Driven Recommendations

R1 — Condition dashboards on the six primary sensors

Action: Visualise **sensor_11, sensor_4, sensor_12, sensor_7, sensor_15, sensor_21** with **20-cycle** rolling means.

Evidence: Figures 5–6.

Expected impact: Earlier degradation visibility; reduced calendar-only dependence.

Ease: Medium | *Priority:* High

R2 — Maintenance triage using a model-estimated risk band

Action: Use the retrospective **RUL ≤ 30** band as the training definition for high-risk states in Assignment 2; in future deployment, trigger inspection when a trained model estimates RUL falls within this range or near-failure probability exceeds a chosen threshold.

Evidence: Figure 3; 15.0% of training rows fall in this band.

Expected impact: Risk-ranked work orders based on model output rather than calendar schedule.

Ease: Low (EDA definition) $\rightarrow$ Medium (deployment) | *Priority:* High

R3 — Deprioritise the seven flat channels

Action: Downsample or archive **sensor_1, sensor_5, sensor_6, sensor_10, sensor_16, sensor_18, sensor_19** at the edge.

Evidence: Figure 1.

Expected impact: Potentially lower ingest and storage cost; limited diagnostic loss in the FD001 training subset, pending validation on other operating conditions.

Ease: Low | *Priority:* Medium

R4 — Engine-aware supervised learning in Assignment 2

Action: Regress **RUL** (and optionally classify `near_failure`) using deployable input features — `time_in_cycles`, top active sensor channels, and rolling sensor summaries — while **holding out entire engines** for validation to prevent temporal leakage.

Evidence: Figures 5, 8, 9; Findings 4–7.

Expected impact: Operational RUL estimates grounded in sensor data available at prediction time.

Ease: Medium | *Priority:* High

R5 — Audit the shortest-lived engines

Action: In a real deployment, cross-reference the **six** units with `max_cycle < 148` cycles against maintenance or inspection logs (not available in the simulation dataset) to determine whether early failure reflects manufacturing variation, unusual operating loads, or data artefacts.

Evidence: Figure 8; `outputs/engine_level_summary.csv`

Expected impact: Operational insight into whether short-life outliers are learnable from sensor history or driven by factors outside the dataset.

Ease: Low (cross-referencing exercise in a real system) | *Priority:* Medium

5.3 Preliminary Modeling Plan for Assignment 2

- **Targets:** RUL regression primary; `near_failure` classification secondary. `capped_RUL` (≤ 125 cycles) may be used as a transformed regression target for modelling stability, but **not as an input feature**.
- **Deployable input features:** `time_in_cycles`, top active sensor channels (`sensor_11`, `sensor_4`, `sensor_12`, `sensor_7`, `sensor_15`, `sensor_21`, `sensor_20`, `sensor_2`, `sensor_17`, `sensor_3`), and their 20-cycle rolling means. Operational settings are retained for completeness. Retrospective variables derived from the failure cycle — including `max_cycle`, `RUL`, `capped_RUL`, `near_failure`, `cycle_ratio`, and `life_stage` — are **not deployable input features**; they are useful for target construction, EDA grouping, and model evaluation only.
- **Validation:** Grouped splits by `unit_number` — entire engines are held out, not individual rows, to prevent temporal data leakage.
- **Models:** Ridge baseline → Random Forest / gradient boosting ensembles; Logistic Regression and Random Forest for classification.
- **Metrics:** RMSE / MAE for regression; F1 / ROC-AUC for alert classification.

6. AI Collaboration Report

Full disclosure: `reports/AI_Collaboration_Log.md`

AI tools assisted with **Python scaffolding**, **plot prototypes**, and **draft wording**; every numeric claim in this report was cross-checked against `outputs/*.csv`, the **executed notebook**, or recomputation (PCA variance). Analytical calls—**retaining outliers**, **seven-sensor low-info policy**, **30-cycle risk band**, and **figure down-selection**—were student-governed per the collaboration log.

7. Continuity Statement for Assignment 2

7.1 Chosen Scenario

Smart Factory Predictive Maintenance — FD001 continuity with `SmartFactory_cleaned.csv`.

7.2 Preliminary Modeling Goals

1. RUL regression
2. `near_failure` classification
3. Optional fleet dashboard tying alerts to rolling telemetry

7.3 Most Promising Features for Assignment 2

The following features are **deployable at prediction time** — they are observable from live sensor streams without knowledge of the engine's future failure cycle.

Feature group	Rationale
<code>time_in_cycles</code>	Available at prediction time; captures accumulated operating age.
Top active sensors: <code>sensor_11</code> , <code>sensor_4</code> , <code>sensor_12</code> , <code>sensor_7</code> , <code>sensor_15</code> , <code>sensor_21</code> , <code>sensor_20</code> , <code>sensor_2</code> , <code>sensor_17</code> , <code>sensor_3</code>	Strong empirical association with retrospective RUL in EDA, supported by Figure 5 and <code>outputs/sensor_rul_correlation.csv</code> .
<code>sensor_*_roll120</code> rolling means for top sensors	Denoised trend features; capture gradual degradation trajectory.
Rolling slopes / deltas (if implemented in notebook)	May represent degradation speed rather than absolute sensor level.
Operational settings (<code>operational_setting_1/2/3</code>)	Low variation in FDOO1; retained for completeness and multi-subset generalisation.

Retrospective variables that must not be used as model inputs:

`max_cycle`, `RUL`, `capped_RUL`, `near_failure`, `cycle_ratio`, `life_stage` — all depend on knowing the future failure cycle. They serve as targets, EDA grouping labels, or evaluation references only.

7.4 Expected Modeling Approach

Ridge baseline → ensembles (RF, GBDT, XGBoost if licensed) with **Purged/grouped CV** across engines plus late-life weighting experimentation if needed.

7.5 Target Variables

Primary `RUL` ; ancillary `near_failure` ; modelling stabiliser `capped_RUL` (≤ 125).

References

- Saxena, A., Goebel, K., Simon, D., & Eklund, N. (2008). *Damage propagation modeling for aircraft engine run-to-failure simulation*. In Proceedings of the 1st International Conference on Prognostics and Health Management (PHMo8), Denver, CO.
- Ramasso, E., & Saxena, A. (2014). *Performance benchmarking and analysis of prognostic methods for CMAPSS datasets*. International Journal of Prognostics and Health Management, 5(2), 1–15.
- NASA Prognostics Center of Excellence (PCoE). *CMAPSS Jet Engine Simulated Data*. Retrieved from <https://www.nasa.gov/intelligent-systems-division/discovery-and-systems-health/pcoe/pcoe-data-set-repository/>

Appendix

All figures (Fig 1–10), cleaned dataset (`data/processed/SmartFactory_cleaned.csv`), output tables (`outputs/*.csv`), and the executed notebook (`notebooks/Assignment1_SmartFactory_EDA.ipynb`) are

included in the project submission folder.